

Set up environment

```
In [1]: import gymnasium as gym  
import numpy as np  
import pandas as pd  
import torch
```

数据准备

```
In [2]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
device
```

```
Out[2]: device(type='cuda')
```

```
In [3]: tech_daily = pd.read_csv(r"data\科技股票.csv")  
tech_daily.set_index('date', inplace=True)  
tech_daily.columns=['AAPL','GOOG','MSFT']  
tech_daily
```

Out[3]:

AAPL GOOG MSFT

date	AAPL	GOOG	MSFT
2022-01-03	178.270326	144.088458	324.504611
2022-01-04	176.007789	143.434930	318.940290
2022-01-05	171.326001	136.717897	306.696845
2022-01-06	168.465997	136.616093	304.273361
2022-01-07	168.632504	136.073308	304.428464
...
2025-11-03	268.789439	284.120000	517.030000
2025-11-04	269.778480	278.060000	514.330000
2025-11-05	269.878383	284.750000	507.160000
2025-11-06	269.508741	285.340000	497.100000
2025-11-07	268.210000	279.700000	496.820000

967 rows × 3 columns

In [4]:

```
debt=pd.read_csv(r"data\无风险.csv",encoding='gbk')
debt.set_index('date', inplace=True)
debt.columns=['US_debt']
debt
```

Out[4]:

US_debt

date	
2022-01-03	1.63
2022-01-04	1.66
2022-01-05	1.71
2022-01-06	1.73
2022-01-07	1.76
...	...
2025-11-03	4.13
2025-11-04	4.10
2025-11-05	4.17
2025-11-06	4.11
2025-11-07	4.11

963 rows × 1 columns

In [5]:

```
tmp=pd.read_csv(r"data\指数和贵金属.csv", encoding='gbk')
tmp.columns=['date','SP500','Gold']
tmp.set_index('date', inplace=True)
tmp
```

Out[5]:

SP500 Gold

date	SP500	Gold
2022-01-03	4796.56	1801.3
2022-01-04	4793.54	1814.9
2022-01-05	4700.58	1810.6
2022-01-06	4696.05	1790.9
2022-01-07	4677.03	1796.5
...
2025-11-03	6851.97	4013.7
2025-11-04	6771.55	3941.3
2025-11-05	6796.29	3990.4
2025-11-06	6720.32	3984.8
2025-11-07	6728.80	4007.8

1001 rows × 2 columns

In [6]:

```
df=pd.merge(tech_daily,debt,how='left',on='date')
df=pd.merge(df,tmp,how='left',on='date')
df['date']=pd.to_datetime(df.index)
df.set_index('date', inplace=True)
df[df.isnull().values == True]
```

Out[6]:

	AAPL	GOOG	MSFT	US_debt	SP500	Gold
date						
2022-10-10	138.103983	98.039090	223.685532	NaN	3612.39	1675.7
2022-11-11	147.475157	96.072547	241.112025	NaN	3992.93	1774.2
2023-10-09	177.082149	138.551849	324.923430	NaN	4335.66	1875.0
2024-10-14	230.005541	165.625113	416.016745	NaN	5859.85	2665.8
2024-11-11	223.220427	181.177047	414.895165	NaN	6001.35	2626.1
2025-10-13	247.420154	244.640000	514.050000	NaN	6654.72	4130.0

In [7]:

```
df.interpolate(method='time', inplace=True)  
df
```

Out[7]:

	AAPL	GOOG	MSFT	US_debt	SP500	Gold
date						
2022-01-03	178.270326	144.088458	324.504611	1.63	4796.56	1801.3
2022-01-04	176.007789	143.434930	318.940290	1.66	4793.54	1814.9
2022-01-05	171.326001	136.717897	306.696845	1.71	4700.58	1810.6
2022-01-06	168.465997	136.616093	304.273361	1.73	4696.05	1790.9
2022-01-07	168.632504	136.073308	304.428464	1.76	4677.03	1796.5
...
2025-11-03	268.789439	284.120000	517.030000	4.13	6851.97	4013.7
2025-11-04	269.778480	278.060000	514.330000	4.10	6771.55	3941.3
2025-11-05	269.878383	284.750000	507.160000	4.17	6796.29	3990.4
2025-11-06	269.508741	285.340000	497.100000	4.11	6720.32	3984.8
2025-11-07	268.210000	279.700000	496.820000	4.11	6728.80	4007.8

967 rows × 6 columns

In [8]:

```
from math import inf

class PortfolioOptimizationEnv(gym.Env):
    def __init__(self, tickers, window_size, start_date, end_date,
                 initial_balance, seed=None):
        super().__init__()
        self.tickers = tickers
        self.window_size = window_size
        self.initial_balance = initial_balance

        # 分别存储原始价格和指标
        self.raw_data, self.feature_data = self.get_data(tickers, start_date, end_date)
        self.n_features = self.feature_data.shape[1]
```

```
self.action_space = gym.spaces.Box(low=0, high=1, shape=(len(tickers),))
self.observation_space = gym.spaces.Box(low=-inf, high=inf,
                                         shape=(window_size, self.n_features))

if seed is not None:
    np.random.seed(seed)
    self.action_space.seed(seed)
    torch.manual_seed(seed)
    if torch.cuda.is_available():
        torch.cuda.manual_seed_all(seed)

def get_data(self, tickers, start_date, end_date):
    data = df.copy().dropna()
    data = data.loc[start_date:end_date, tickers]

    # 保存原始价格（用于计算投资组合收益）
    raw_data = data.copy()

    # 计算特征指标
    returns = data.pct_change()

    mom_frames = []
    for window in [5, 20]:
        mom = data / data.shift(window) - 1
        mom.columns = [f"{col}_mom_{window}" for col in data.columns]
        mom_frames.append(mom)

    vol = returns.rolling(window=20, min_periods=1).std()
    vol.columns = [f"{col}_vol_20" for col in data.columns]

    ma = data.rolling(window=20, min_periods=1).mean()
    ma_dev = data / ma - 1
    ma_dev.columns = [f"{col}_ma_dev_20" for col in data.columns]

    returns.columns = [f"{col}_ret" for col in data.columns]

    # 特征数据: returns, vol, ma_dev, momentum (不包含原始价格)
    feature_data = pd.concat([returns, vol, ma_dev] + mom_frames, axis=1)
    raw_data = raw_data.dropna()
    feature_data = feature_data.reindex(raw_data.index)
    feature_data.fillna(method='ffill', inplace=True)
```

```
feature_data.fillna(method='bfill', inplace=True)

return raw_data.dropna(), feature_data.dropna()

def reset(self, seed=None):
    self.balance = self.initial_balance
    self.current_step = self.window_size

    # 使用特征数据作为观察
    obs = self.feature_data.iloc[self.current_step - self.window_size:self.current_step].values
    info = {"balance": self.balance}
    return obs, info

def step(self, action):
    action = action / (np.sum(action) + 1e-8)

    prev_balance = self.balance

    # 从原始价格计算实际收益
    current_price = self.raw_data.iloc[self.current_step].values[:len(self.tickers)]
    prev_price = self.raw_data.iloc[self.current_step - 1].values[:len(self.tickers)]
    asset_returns = current_price / prev_price - 1

    # 基础奖励：投资组合收益
    portfolio_return = np.sum(asset_returns * action)
    self.balance = self.balance * (1 + portfolio_return)
    base_reward = np.log(self.balance / prev_balance)

    # 可选：添加风险惩罚
    # 从特征数据中获取波动率
    current_features = self.feature_data.iloc[self.current_step].values
    volatility = current_features[len(self.tickers):2*len(self.tickers)].mean() # 平均波动率
    risk_penalty = -0.1 * volatility

    # 总奖励
    reward = base_reward + risk_penalty

    self.current_step += 1
    done = self.current_step >= len(self.raw_data)-1

    obs_end=min(len(self.feature_data),self.current_step+self.window_size)
```

```
        obs_start=max(0,obs_end - self.window_size)
        obs = self.feature_data.iloc[obs_start:obs_end].values

        terminated = bool(done)
        truncated = False
        info = {'balance': self.balance}

    return obs, reward, terminated, truncated, info
```

```
In [9]: tickers = df.columns.tolist()
window_size = 30
start_date = '2022-01-01'
end_date = '2025-09-01'
initial_balance = 10000
seed = 8

# Initialize the environment
env = PortfolioOptimizationEnv(
    tickers,
    window_size,
    start_date,
    end_date,
    initial_balance,
    seed)

# Get the initial state
state = env.reset(seed=seed)
# Sample and execute a random action
action = env.action_space.sample()
next_state, reward, terminated, truncated, info = env.step(action)
done = bool(terminated or truncated)
# print(f"State: {state}")
print(f"Action: {action}")
# print(f"Next state: {next_state}")
print(f"Reward: {reward}")
print(f"Balance: {info['balance']}")
print(f"Done: {done}")
```

```
Action: [0.32697228 0.98727685 0.31871083 0.78854895 0.86989653 0.39108482]
```

```
Reward: 0.014046373702849141
```

```
Balance: 10161.65097308565
```

```
Done: False
```

```
C:\Users\HP\AppData\Local\Temp\ipykernel_49456\356491190.py:55: FutureWarning: DataFrame.fillna with 'method' is deprecated and  
will raise in a future version. Use obj.ffill() or obj.bfill() instead.
```

```
    feature_data.fillna(method='ffill', inplace=True)
```

```
C:\Users\HP\AppData\Local\Temp\ipykernel_49456\356491190.py:56: FutureWarning: DataFrame.fillna with 'method' is deprecated and  
will raise in a future version. Use obj.ffill() or obj.bfill() instead.
```

```
    feature_data.fillna(method='bfill', inplace=True)
```

Training DRL agent

```
In [10]: # ===== Train PPO + plot training profile + plot cumulative wealth + print metrics =====  
# Prereqs (run once in your env if needed):  
#   pip install "stable-baselines3>=2.3.0" "shimmy>=2.0" matplotlib pandas  
  
from stable_baselines3 import PPO  
from stable_baselines3.common.vec_env import DummyVecEnv  
from stable_baselines3.common.monitor import Monitor  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import os  
  
# ----- 1) Build vectorized env (wrap with Monitor to log training rewards) -----  
log_dir = "./sb3_logs"  
os.makedirs(log_dir, exist_ok=True)  
monitor_path = os.path.join(log_dir, "monitor.csv")  
  
def make_env():  
    env_ = PortfolioOptimizationEnv(  
        tickers=tickers,  
        window_size=window_size,  
        start_date=start_date,  
        end_date=end_date,  
        initial_balance=initial_balance,  
        seed=seed,
```

```

)
# Log episode reward/length; include "balance" from info if you want (optional)
env_ = Monitor(env_, filename=monitor_path) # single-env -> single monitor file
return env_

vec_env = DummyVecEnv([make_env])

# ----- 2) Train PPO -----
model = PPO(
    "MlpPolicy", vec_env, verbose=1, device=device
) # PPO (近端策略优化)：适合连续动作空间，训练稳定，是投资组合优化的常用算法
model.learn(total_timesteps=20000) # increase as needed
model.save("ppo_portfolio_optimization")

# ----- 3) Load training profile (episode rewards) and plot -----
# SB3 Monitor CSV starts with commented metadata lines beginning with '#'
train_df = pd.read_csv(monitor_path, comment="#")
# Columns typically: r (ep reward), l (ep length), t (time)
# Make a simple moving average of episode rewards for a smooth training curve
if len(train_df) > 0:
    train_df["ep"] = np.arange(1, len(train_df) + 1)
    train_df["reward_smooth"] = train_df["r"].rolling(window=max(5, len(train_df)//50), min_periods=1).mean()

    plt.figure(figsize=(8, 4))
    plt.plot(train_df["ep"], train_df["r"], alpha=0.3, label="Episode reward")
    plt.plot(train_df["ep"], train_df["reward_smooth"], label="Smoothed")
    plt.xlabel("Episode")
    plt.ylabel("Episode reward")
    plt.title("Training profile (episode rewards)")
    plt.legend()
    plt.tight_layout()
    plt.show()
else:
    print("No episodes logged in monitor file (did training terminate too early?).")

# ----- 4) Roll out a full episode deterministically and record balances -----
obs = vec_env.reset()
dones = [False]
balances = []
rewards = []
weights_seq = []

```

```

while not dones[0]:
    action, _ = model.predict(obs, deterministic=True)
    obs, reward, done, infos = vec_env.step(action)
    rewards.append(float(reward[0]))
    balances.append(float(infos[0].get("balance", np.nan)))
    weights_seq.append(infos[0].get("weights", None))

# Try to build a date index for nicer plotting/annualization
try:
    # Access the underlying single env inside DummyVecEnv -> Monitor -> PortfolioOptimizationEnv
    base_env = vec_env.envs[0].env
    # For safety if nested wrappers differ:
    while hasattr(base_env, "env"):
        base_env = base_env.env
    price_index = base_env.data.index # pandas.DatetimeIndex from your env
    # The first balance corresponds to the first step after initial window
    # Align dates accordingly:
    start_idx = window_size
    end_idx = start_idx + len(balances)
    dates = price_index[start_idx:end_idx]
except Exception:
    dates = pd.RangeIndex(start=1, stop=len(balances) + 1) # fallback to simple index

# ----- 5) Plot cumulative wealth curve -----
plt.figure(figsize=(8, 4))
plt.plot(dates, balances)
plt.xlabel("Date" if isinstance(dates, pd.DatetimeIndex) else "Step")
plt.ylabel("Wealth")
plt.title("Cumulative wealth (deterministic policy rollout)")
plt.tight_layout()
plt.show()

# ----- 6) Compute portfolio performance metrics -----
# Step simple returns from balance path
balances_arr = np.asarray(balances, dtype=float)
rets = balances_arr[1:] / balances_arr[:-1] - 1.0 if len(balances_arr) > 1 else np.array([])

# Infer periods per year from dates if possible; else default 252
def infer_ppy(idx):
    if isinstance(idx, pd.DatetimeIndex) and len(idx) >= 2:

```

```

# Use median days between points
deltas = np.diff(idx.view("int64")) / 1e9 / 86400.0 # ns -> days
median_days = np.median(deltas) if len(deltas) else 1.0
if median_days <= 0:
    return 252.0
return 365.25 / median_days
return 252.0

PPY = infer_ppy(dates)

def max_drawdown(equity):
    peak = np.maximum.accumulate(equity)
    dd = (equity - peak) / peak
    return float(dd.min()) if len(dd) else 0.0

if len(rets) > 0:
    total_growth = balances_arr[-1] / balances_arr[0]
    N = len(rets)
    CAGR = total_growth ** (PPY / N) - 1.0 # PPY: 每年交易天数; CAGR: 年复合增长率
    mean_r = float(np.mean(rets))
    std_r = float(np.std(rets, ddof=1)) if N > 1 else 0.0
    ann_return = (1.0 + mean_r) ** PPY - 1.0 if mean_r > -1 else np.nan # geometric approx alternative below
    ann_vol = std_r * np.sqrt(PPY) if std_r > 0 else 0.0
    # Sharpe (rf=0 for simplicity; replace with your risk-free rate if desired)
    rf = 0.0
    sharpe = (mean_r - rf/PPY) / std_r * np.sqrt(PPY) if std_r > 0 else np.nan
    mdd = max_drawdown(balances_arr)
    calmar = (CAGR / abs(mdd)) if mdd != 0 else np.nan

    # Alternative geometric annualized return from total growth (more stable):
    ann_return_geom = total_growth ** (PPY / N) - 1.0

print("\n==== Portfolio Performance (deterministic rollout) ===")
print(f"Periods per year (inferred): {PPY:.2f}")
print(f"Start wealth: {balances_arr[0]:.2f}")
print(f"End wealth: {balances_arr[-1]:.2f}")
print(f"Total growth: {total_growth:.6f}")
print(f"CAGR: {CAGR:.6%}")
print(f"Ann. Return (geom from equity): {ann_return_geom:.6%}")
print(f"Ann. Vol: {ann_vol:.6%}")
print(f"Sharpe (rf=0): {sharpe:.4f}")

```

```

        print(f"Max Drawdown: {mdd:.2%}")
        print(f"Calmar: {calmar:.4f}")
else:
    print("Not enough steps to compute metrics (need at least 2 balances.)")

# ----- 7) (Optional) Inspect last action weights -----
if any(w is not None for w in weights_seq):
    last_w = [w for w in weights_seq if w is not None][-1]
    print("\nLast portfolio weights:", np.round(last_w, 4))

```

Gym has been unmaintained since 2022 and does not support NumPy 2.0 amongst other critical functionality.
Please upgrade to Gymnasium, the maintained drop-in replacement of Gym, or contact the authors of your software and request that they upgrade.

Users of this version of Gym should be able to simply replace 'import gym' with 'import gymnasium as gym' in the vast majority of cases.

See the migration guide at https://gymnasium.farama.org/introduction/migration_guide/ for additional information.

C:\Users\HP\AppData\Local\Temp\ipykernel_49456\356491190.py:55: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.

feature_data.fillna(method='ffill', inplace=True)

C:\Users\HP\AppData\Local\Temp\ipykernel_49456\356491190.py:56: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.

feature_data.fillna(method='bfill', inplace=True)

Using cuda device

d:\Develop\miniconda\envs\py310\lib\site-packages\stable_baselines3\common\on_policy_algorithm.py:150: UserWarning: You are trying to run PPO on the GPU, but it is primarily intended to run on the CPU when not using a CNN policy (you are using ActorCriticPolicy which should be a MlpPolicy). See <https://github.com/DLR-RM/stable-baselines3/issues/1245> for more info. You can pass `device='cpu'` or `export CUDA_VISIBLE_DEVICES=` to force using the CPU. Note: The model will train, but the GPU utilization will be poor and the training might take longer than on CPU.

warnings.warn(

rollout/		
ep_len_mean	887	
ep_rew_mean	-1.21	
time/		
fps	462	
iterations	1	
time_elapsed	4	
total_timesteps	2048	

rollout/		
ep_len_mean	887	
ep_rew_mean	-1.1	
time/		
fps	389	
iterations	2	
time_elapsed	10	
total_timesteps	4096	
train/		
approx_kl	0.011333423	
clip_fraction	0.078	
clip_range	0.2	
entropy_loss	-8.52	
explained_variance	-0.417	
learning_rate	0.0003	
loss	-0.0559	
n_updates	10	
policy_gradient_loss	-0.018	
std	1	
value_loss	0.000941	

rollout/		
ep_len_mean	887	
ep_rew_mean	-0.902	
time/		
fps	350	
iterations	3	
time_elapsed	17	
total_timesteps	6144	

train/		
approx_kl	0.012782114	
clip_fraction	0.109	
clip_range	0.2	
entropy_loss	-8.52	
explained_variance	0.559	
learning_rate	0.0003	
loss	-0.0598	
n_updates	20	
policy_gradient_loss	-0.0254	
std	1	
value_loss	0.000681	

rollout/		
ep_len_mean	887	
ep_rew_mean	-0.772	
time/		
fps	328	
iterations	4	
time_elapsed	24	
total_timesteps	8192	
train/		
approx_kl	0.014021457	
clip_fraction	0.135	
clip_range	0.2	
entropy_loss	-8.52	
explained_variance	0.617	
learning_rate	0.0003	
loss	-0.0407	
n_updates	30	
policy_gradient_loss	-0.0293	
std	1	
value_loss	0.000633	

rollout/		
ep_len_mean	887	
ep_rew_mean	-0.639	
time/		
fps	317	

iterations	5
time_elapsed	32
total_timesteps	10240
<hr/>	
train/	
approx_kl	0.013396604
clip_fraction	0.124
clip_range	0.2
entropy_loss	-8.49
explained_variance	0.693
learning_rate	0.0003
loss	-0.0467
n_updates	40
policy_gradient_loss	-0.0275
std	0.993
value_loss	0.000668
<hr/>	
rollout/	
ep_len_mean	887
ep_rew_mean	-0.514
time/	
fps	309
iterations	6
time_elapsed	39
total_timesteps	12288
train/	
approx_kl	0.016518926
clip_fraction	0.152
clip_range	0.2
entropy_loss	-8.48
explained_variance	0.688
learning_rate	0.0003
loss	-0.0449
n_updates	50
policy_gradient_loss	-0.0321
std	0.995
value_loss	0.000507
<hr/>	
rollout/	
ep_len_mean	887

ep_rew_mean	-0.37
time/	
fps	305
iterations	7
time_elapsed	46
total_timesteps	14336
train/	
approx_kl	0.01791879
clip_fraction	0.171
clip_range	0.2
entropy_loss	-8.46
explained_variance	0.65
learning_rate	0.0003
loss	-0.047
n_updates	60
policy_gradient_loss	-0.0354
std	0.987
value_loss	0.000468


rollout/	
ep_len_mean	887
ep_rew_mean	-0.254
time/	
fps	300
iterations	8
time_elapsed	54
total_timesteps	16384
train/	
approx_kl	0.019249577
clip_fraction	0.184
clip_range	0.2
entropy_loss	-8.42
explained_variance	0.668
learning_rate	0.0003
loss	-0.0549
n_updates	70
policy_gradient_loss	-0.0344
std	0.984
value_loss	0.000515

rollout/	
ep_len_mean	887
ep_rew_mean	-0.116
time/	
fps	298
iterations	9
time_elapsed	61
total_timesteps	18432
train/	
approx_kl	0.020216336
clip_fraction	0.194
clip_range	0.2
entropy_loss	-8.38
explained_variance	0.573
learning_rate	0.0003
loss	-0.0614
n_updates	80
policy_gradient_loss	-0.0373
std	0.974
value_loss	0.000498

rollout/	
ep_len_mean	887
ep_rew_mean	0.0532
time/	
fps	296
iterations	10
time_elapsed	68
total_timesteps	20480
train/	
approx_kl	0.018311683
clip_fraction	0.193
clip_range	0.2
entropy_loss	-8.32
explained_variance	0.638
learning_rate	0.0003
loss	-0.0661
n_updates	90
policy_gradient_loss	-0.0364

std	0.964
value_loss	0.000426





==== Portfolio Performance (deterministic rollout) ===

Periods per year (inferred): 252.00

Start wealth: 10297.47

End wealth: 23731074.08

Total growth: 2304.553868

CAGR: 804.485651%

Ann. Return (geom from equity): 804.485651%

Ann. Vol: 22.710210%

Sharpe (rf=0): 9.8502

Max Drawdown: -5.59%

Calmar: 143.9561

修改部分

添加了科技股票、国债、标普500和黄金等数据

