

# Set up environment

```
In [21]: import gymnasium as gym  
import numpy as np  
import pandas as pd  
import torch
```

## 数据准备

```
In [22]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
device
```

```
Out[22]: device(type='cuda')
```

```
In [23]: tech_daily = pd.read_csv(r"data\科技股票.csv")  
tech_daily.set_index('date', inplace=True)  
tech_daily.columns=['AAPL','GOOG','MSFT']  
tech_daily
```

Out[23]:

	AAPL	GOOG	MSFT
date			
2022-01-03	178.270326	144.088458	324.504611
2022-01-04	176.007789	143.434930	318.940290
2022-01-05	171.326001	136.717897	306.696845
2022-01-06	168.465997	136.616093	304.273361
2022-01-07	168.632504	136.073308	304.428464
...	...	...	...
2025-11-03	268.789439	284.120000	517.030000
2025-11-04	269.778480	278.060000	514.330000
2025-11-05	269.878383	284.750000	507.160000
2025-11-06	269.508741	285.340000	497.100000
2025-11-07	268.210000	279.700000	496.820000

967 rows × 3 columns

In [24]:

```
debt=pd.read_csv(r"data\无风险.csv",encoding='gbk')
debt.set_index('date', inplace=True)
debt.columns=['US_debt']
debt
```

Out[24]:

**US\_debt**

date	
2022-01-03	1.63
2022-01-04	1.66
2022-01-05	1.71
2022-01-06	1.73
2022-01-07	1.76
...	...
2025-11-03	4.13
2025-11-04	4.10
2025-11-05	4.17
2025-11-06	4.11
2025-11-07	4.11

963 rows × 1 columns

In [25]:

```
tmp=pd.read_csv(r"data\指数和贵金属.csv", encoding='gbk')
tmp.columns=['date','SP500','Gold']
tmp.set_index('date', inplace=True)
tmp
```

Out[25]:

SP500 Gold

date	SP500	Gold
2022-01-03	4796.56	1801.3
2022-01-04	4793.54	1814.9
2022-01-05	4700.58	1810.6
2022-01-06	4696.05	1790.9
2022-01-07	4677.03	1796.5
...	...	...
2025-11-03	6851.97	4013.7
2025-11-04	6771.55	3941.3
2025-11-05	6796.29	3990.4
2025-11-06	6720.32	3984.8
2025-11-07	6728.80	4007.8

1001 rows × 2 columns

In [26]:

```
df=pd.merge(tech_daily,debt,how='left',on='date')
df=pd.merge(df,tmp,how='left',on='date')
df['date']=pd.to_datetime(df.index)
df.set_index('date', inplace=True)
df[df.isnull().values == True]
```

Out[26]:

	AAPL	GOOG	MSFT	US_debt	SP500	Gold
date						
2022-10-10	138.103983	98.039090	223.685532	NaN	3612.39	1675.7
2022-11-11	147.475157	96.072547	241.112025	NaN	3992.93	1774.2
2023-10-09	177.082149	138.551849	324.923430	NaN	4335.66	1875.0
2024-10-14	230.005541	165.625113	416.016745	NaN	5859.85	2665.8
2024-11-11	223.220427	181.177047	414.895165	NaN	6001.35	2626.1
2025-10-13	247.420154	244.640000	514.050000	NaN	6654.72	4130.0

In [27]:

```
df.interpolate(method='time', inplace=True)  
df
```

Out[27]:

	AAPL	GOOG	MSFT	US_debt	SP500	Gold
date						
2022-01-03	178.270326	144.088458	324.504611	1.63	4796.56	1801.3
2022-01-04	176.007789	143.434930	318.940290	1.66	4793.54	1814.9
2022-01-05	171.326001	136.717897	306.696845	1.71	4700.58	1810.6
2022-01-06	168.465997	136.616093	304.273361	1.73	4696.05	1790.9
2022-01-07	168.632504	136.073308	304.428464	1.76	4677.03	1796.5
...	...	...	...	...	...	...
2025-11-03	268.789439	284.120000	517.030000	4.13	6851.97	4013.7
2025-11-04	269.778480	278.060000	514.330000	4.10	6771.55	3941.3
2025-11-05	269.878383	284.750000	507.160000	4.17	6796.29	3990.4
2025-11-06	269.508741	285.340000	497.100000	4.11	6720.32	3984.8
2025-11-07	268.210000	279.700000	496.820000	4.11	6728.80	4007.8

967 rows × 6 columns

In [28]:

```
# 保证日期有序
df = df.sort_index()

print("数据日期范围:", df.index.min(), "→", df.index.max())

# 生成滑动窗口函数
def generate_walk_forward_windows(index, train_years=2, test_months=6, step_months=6):
    """
    基于日期索引生成滑动窗口：
    - 每个窗口：训练 train_years 年，测试 test_months 个月
    - 每次向前滚动 step_months 个月
    返回：
        [{"train_start":..., "train_end":..., "test_start":..., "test_end":...}, ...]
    """

```

```
"""
windows = []

start_date = index.min().normalize()
last_date = index.max().normalize()

current_train_start = start_date

while True:
    train_start = current_train_start
    train_end = (
        train_start + pd.DateOffset(years=train_years) - pd.DateOffset(days=1)
    )
    test_start = train_end + pd.DateOffset(days=1)
    test_end = (
        test_start + pd.DateOffset(months=test_months) - pd.DateOffset(days=1)
    )

    # 测试区间超出数据范围就停止
    if test_end > last_date:
        break

    windows.append(
    {
        "train_start": train_start,
        "train_end": train_end,
        "test_start": test_start,
        "test_end": test_end,
    }
)

# 向前滚动 step_months 个月
current_train_start = current_train_start + pd.DateOffset(months=step_months)

return windows

# 生成窗口（你可以根据需要改 train_years / test_months / step_months）
windows = generate_walk_forward_windows(
    df.index, train_years=2, test_months=6, step_months=6
)
```

```
print(f"共生成 {len(windows)} 个滑动窗口: ")
for i, w in enumerate(windows):
    print(
        f"Window {i}: "
        f"Train {w['train_start'].date()} ~ {w['train_end'].date()} | "
        f"Test {w['test_start'].date()} ~ {w['test_end'].date()}"
    )
```

数据日期范围: 2022-01-03 00:00:00 → 2025-11-07 00:00:00

共生成 3 个滑动窗口:

```
Window 0: Train 2022-01-03 ~ 2024-01-02 | Test 2024-01-03 ~ 2024-07-02
Window 1: Train 2022-07-03 ~ 2024-07-02 | Test 2024-07-03 ~ 2025-01-02
Window 2: Train 2023-01-03 ~ 2025-01-02 | Test 2025-01-03 ~ 2025-07-02
```

In [29]:

```
from math import inf

from networkx import sigma
from pyparsing import deque

class PortfolioOptimizationEnv(gym.Env):
    def __init__(self, tickers, window_size, start_date, end_date, initial_balance, seed=None):
        super().__init__()
        self.tickers = tickers
        self.window_size = window_size
        self.initial_balance = initial_balance

        # 分别存储原始价格和指标
        self.raw_data, self.feature_data = self.get_data(tickers, start_date, end_date)
        self.n_features = self.feature_data.shape[1]

        self.action_space = gym.spaces.Box(low=0, high=1, shape=(len(tickers),))
        self.observation_space = gym.spaces.Box(
            low=-inf, high=inf, shape=(window_size, self.n_features)
        )

        self.return_window = deque(maxlen=window_size)
        self.last_action = np.ones(len(tickers)) / len(tickers)
```

```
if seed is not None:
    np.random.seed(seed)
    self.action_space.seed(seed)
    torch.manual_seed(seed)
    if torch.cuda.is_available():
        torch.cuda.manual_seed_all(seed)

def get_data(self, tickers, start_date, end_date):
    data = df.copy().dropna()
    data = data.loc[start_date:end_date, tickers]

    # 保存原始价格（用于计算投资组合收益）
    raw_data = data.copy()

    # 计算特征指标
    returns = data.pct_change()

    mom_frames = []
    for window in [5, 20]:
        mom = data / data.shift(window) - 1
        mom.columns = [f"{col}_mom_{window}" for col in data.columns]
        mom_frames.append(mom)

    vol = returns.rolling(window=20, min_periods=1).std()
    vol.columns = [f"{col}_vol_20" for col in data.columns]

    ma = data.rolling(window=20, min_periods=1).mean()
    ma_dev = data / ma - 1
    ma_dev.columns = [f"{col}_ma_dev_20" for col in data.columns]

    returns.columns = [f"{col}_ret" for col in data.columns]

    # 特征数据: returns, vol, ma_dev, momentum (不包含原始价格)
    feature_data = pd.concat([returns, vol, ma_dev] + mom_frames, axis=1)
    raw_data = raw_data.dropna()
    feature_data = feature_data.reindex(raw_data.index)
    feature_data.fillna(method="ffill", inplace=True)
    feature_data.fillna(method="bfill", inplace=True)

return raw_data.dropna(), feature_data.dropna()
```

```
def reset(self, seed=None):
    self.balance = self.initial_balance
    self.current_step = self.window_size

    self.return_window.clear()
    self.last_action = np.ones(len(self.tickers)) / len(self.tickers)

    # 使用特征数据作为观察
    obs = self.feature_data.iloc[
        self.current_step - self.window_size : self.current_step
    ].values
    info = {"balance": self.balance}
    return obs, info

def step(self, action):
    # SAC已经输出归一化的动作，因此不需要进行softmax
    action = np.asarray(action).ravel()
    action = np.clip(action, 0, 1)
    action = action / np.sum(action + 1e-8)

    prev_balance = self.balance

    # 从原始价格计算实际收益
    current_price = self.raw_data.iloc[self.current_step].values[
        : len(self.tickers)
    ]
    prev_price = self.raw_data.iloc[self.current_step - 1].values[
        : len(self.tickers)
    ]
    asset_returns = current_price / prev_price - 1

    self.return_window.append(asset_returns)

    # 基础奖励：投资组合收益
    portfolio_return = np.sum(asset_returns * action)
    self.balance = self.balance * (1 + portfolio_return)
    base_reward = np.log(self.balance / prev_balance)

    risk_penalty = 0
    if len(self.return_window) >= 5:
```

```

        R = np.vstack(self.return_window)
        cov_matrix = np.cov(R.T)
        sigma_p2 = action.T @ cov_matrix @ action
        risk_penalty = sigma_p2

        turnover = np.sum(np.abs(action - self.last_action))
        cost = turnover
        self.last_action = action

        # 总奖励
        lambda_risk = 1
        lambda_turnover = 0.001
        reward = base_reward - lambda_risk * risk_penalty - lambda_turnover * cost

        self.current_step += 1
        done = self.current_step >= len(self.raw_data) - 1

        obs_end = min(len(self.feature_data), self.current_step + self.window_size)
        obs_start = max(0, obs_end - self.window_size)
        obs = self.feature_data.iloc[obs_start:obs_end].values

        terminated = bool(done)
        truncated = False
        info = {"balance": self.balance}

    return obs, reward, terminated, truncated, info

```

In [30]:

```

tickers = df.columns.tolist()
window_size = 30
start_date = '2022-01-01'
end_date = '2025-09-01'
initial_balance = 10000
seed = 8

# Initialize the environment
env = PortfolioOptimizationEnv(
    tickers,
    window_size,
    start_date,
    end_date,

```

```

    initial_balance,
    seed)

# Get the initial state
state = env.reset(seed=seed)
# Sample and execute a random action
action = env.action_space.sample()
next_state, reward, terminated, truncated, info = env.step(action)
done = bool(terminated or truncated)
# print(f"State: {state}")
print(f"Action: {action}")
# print(f"Next state: {next_state}")
print(f"Reward: {reward}")
print(f"Balance: {info['balance']}")
print(f"Done: {done}")

```

Action: [0.32697228 0.98727685 0.31871083 0.78854895 0.86989653 0.39108482]  
Reward: 0.015598913072904633  
Balance: 10161.65097308565  
Done: False

C:\Users\HP\AppData\Local\Temp\ipykernel\_9704\1757598842.py:64: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.  
feature\_data.fillna(method="ffill", inplace=True)  
C:\Users\HP\AppData\Local\Temp\ipykernel\_9704\1757598842.py:65: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.  
feature\_data.fillna(method="bfill", inplace=True)

## Training DRL agent

### SAC

In [31]:

```
# ===== Walk-Forward 滑动窗口训练 + 测试 SAC =====
from stable_baselines3 import SAC
from stable_baselines3.common.vec_env import DummyVecEnv, VecNormalize
from stable_baselines3.common.monitor import Monitor
import numpy as np
import pandas as pd
```

```
import matplotlib.pyplot as plt
import os

# 日志目录
log_dir = "./sb3_logs_wf"
os.makedirs(log_dir, exist_ok=True)

# 为某个时间段构建一个 env 的工厂函数
def make_env_for_period(start_date_str, end_date_str, monitor_file=None):
    """
    start_date_str, end_date_str: 'YYYY-MM-DD'
    monitor_file: 如不为 None, 则用 Monitor 记录训练信息
    """

    def __init__():
        env_ = PortfolioOptimizationEnv(
            tickers=tickers,
            window_size=window_size,
            start_date=start_date_str,
            end_date=end_date_str,
            initial_balance=initial_balance,
            seed=seed,
        )
        if monitor_file is not None:
            env_ = Monitor(env_, filename=monitor_file)
        return env_

    return __init__

# 保存每个窗口的测试结果
wf_results = []

# 训练总步数（你可以调大, 比如 50_000 / 100_000）
TOTAL_TIMESTEPS = 20_000

for i, w in enumerate(windows):
    print("=" * 80)
    print(f"Window {i}")
    print(f"Train: {w['train_start'].date()} ~ {w['train_end'].date()}"
```

```
print(f"Test : {w['test_start'].date()} ~ {w['test_end'].date()}")  
  
train_start_str = w["train_start"].strftime("%Y-%m-%d")  
train_end_str = w["train_end"].strftime("%Y-%m-%d")  
test_start_str = w["test_start"].strftime("%Y-%m-%d")  
test_end_str = w["test_end"].strftime("%Y-%m-%d")  
  
# ----- 1) 构建训练环境 + 归一化 -----  
monitor_path = os.path.join(log_dir, f"monitor_train_window_{i}.csv")  
  
vec_train_env = DummyVecEnv(  
    [make_env_for_period(train_start_str, train_end_str, monitor_file=monitor_path)])  
)  
vec_train_env = VecNormalize(  
    vec_train_env, norm_obs=True, norm_reward=False, clip_obs=10.0  
)  
  
# ----- 2) 训练 SAC -----  
model = SAC(  
    "MlpPolicy",  
    vec_train_env,  
    verbose=1,  
    device=device, # 你在前面 Cell 3 已定义 device  
    # 下面这些超参数可以之后再细调  
    learning_rate=3e-4,  
    batch_size=256,  
    buffer_size=200_000,  
    tau=0.005,  
    ent_coef="auto",  
    train_freq=(1, "step"),  
    gradient_steps=1,  
)  
  
model.learn(total_timesteps=TOTAL_TIMESTEPS)  
  
# 保存模型和归一化器（方便以后单独加载某个窗口）  
model_path = os.path.join(log_dir, f"sac_window_{i}.zip")  
vecnorm_path = os.path.join(log_dir, f"sac_vecnorm_window_{i}.pkl")  
model.save(model_path)  
vec_train_env.save(vecnorm_path)  
print(f"Saved model to {model_path}")
```

```
print(f"Saved VecNormalize to {vecnorm_path}")

# ----- 3) 构建测试环境（使用同一归一化统计） -----
test_env_raw = DummyVecEnv(
    [make_env_for_period(test_start_str, test_end_str, monitor_file=None)])
)
test_env = VecNormalize.load(vecnorm_path, test_env_raw)
test_env.training = False
test_env.norm_reward = False # 测试时不要再归一化 reward

# ----- 4) 在测试集上回测该窗口策略 -----
obs = test_env.reset()
dones = [False]
balances = []
rewards = []
weights_seq = []

while not done[0]:
    action, _ = model.predict(obs, deterministic=True)
    obs, reward, done, infos = test_env.step(action)
    rewards.append(float(reward[0]))
    balances.append(float(infos[0].get("balance", np.nan)))
    weights_seq.append(infos[0].get("weights", None))

# 尝试从测试环境中拿到日期索引
try:
    base_env = test_env.envs[0].env
    while hasattr(base_env, "env"):
        base_env = base_env.env
    price_index = base_env.raw_data.index # 你环境里存的是 raw_data
    start_idx = window_size
    end_idx = start_idx + len(balances)
    dates = price_index[start_idx:end_idx]
except Exception:
    dates = pd.RangeIndex(start=1, stop=len(balances) + 1)

balances_arr = np.asarray(balances, dtype=float)
rets = (
    balances_arr[1:] / balances_arr[:-1] - 1.0
    if len(balances_arr) > 1
    else np.array([])
```

```
)  
  
# 简单算一下每个窗口的指标（年化收益、波动、Sharpe）  
def infer_ppy(idx):  
    if isinstance(idx, pd.DatetimeIndex) and len(idx) >= 2:  
        deltas = np.diff(idx.view("int64")) / 1e9 / 86400.0  
        median_days = np.median(deltas) if len(deltas) else 1.0  
        if median_days <= 0:  
            return 252.0  
        return 365.25 / median_days  
    return 252.0  
  
PPY = infer_ppy(dates)  
  
if len(rets) > 1:  
    total_growth = balances_arr[-1] / balances_arr[0]  
    ann_return_geom = total_growth ** (PPY / len(rets)) - 1.0  
    ann_vol = np.std(rets) * np.sqrt(PPY)  
    sharpe = ann_return_geom / ann_vol if ann_vol > 1e-12 else np.nan  
else:  
    total_growth = np.nan  
    ann_return_geom = np.nan  
    ann_vol = np.nan  
    sharpe = np.nan  
  
wf_results.append(  
{  
    "window_id": i,  
    "train_start": w["train_start"],  
    "train_end": w["train_end"],  
    "test_start": w["test_start"],  
    "test_end": w["test_end"],  
    "test_steps": len(rets),  
    "start_wealth": balances_arr[0] if len(balances_arr) > 0 else np.nan,  
    "end_wealth": balances_arr[-1] if len(balances_arr) > 0 else np.nan,  
    "total_growth": total_growth,  
    "ann_return": ann_return_geom,  
    "ann_vol": ann_vol,  
    "sharpe": sharpe,  
}  
)
```

```
print(f"Window {i} 测试结束: ")
print(f" 终值资产: {balances_arr[-1]:.2f} (初始 {balances_arr[0]:.2f})")
print(f" 总增长倍数: {total_growth:.4f}")
print(f" 年化收益: {ann_return_geom:.2%}")
print(f" 年化波动: {ann_vol:.2%}")
print(f" Sharpe: {sharpe:.3f}")

# ----- 5) 汇总所有窗口结果 -----
wf_results_df = pd.DataFrame(wf_results)
wf_results_df
```

```
=====
Window 0
Train: 2022-01-03 ~ 2024-01-02
Test : 2024-01-03 ~ 2024-07-02
Using cuda device
C:\Users\HP\AppData\Local\Temp\ipykernel_9704\1757598842.py:64: FutureWarning: DataFrame.fillna with 'method' is deprecated and
will raise in a future version. Use obj.ffill() or obj.bfill() instead.
    feature_data.fillna(method="ffill", inplace=True)
C:\Users\HP\AppData\Local\Temp\ipykernel_9704\1757598842.py:65: FutureWarning: DataFrame.fillna with 'method' is deprecated and
will raise in a future version. Use obj.ffill() or obj.bfill() instead.
    feature_data.fillna(method="bfill", inplace=True)
```

rollout/		
ep_len_mean	471	
ep_rew_mean	-0.0944	
time/		
episodes	4	
fps	42	
time_elapsed	43	
total_timesteps	1884	
train/		
actor_loss	-26.7	
critic_loss	1.44	
ent_coef	0.586	
ent_coef_loss	-5.38	
learning_rate	0.0003	
n_updates	1783	

rollout/		
ep_len_mean	471	
ep_rew_mean	-0.0647	
time/		
episodes	8	
fps	42	
time_elapsed	87	
total_timesteps	3768	
train/		
actor_loss	-38.3	
critic_loss	0.37	
ent_coef	0.333	
ent_coef_loss	-11	
learning_rate	0.0003	
n_updates	3667	

rollout/		
ep_len_mean	471	
ep_rew_mean	-0.0934	
time/		
episodes	12	
fps	43	

time_elapsed	130
total_timesteps	5652
train/	
actor_loss	-41.8
critic_loss	0.048
ent_coef	0.19
ent_coef_loss	-16.6
learning_rate	0.0003
n_updates	5551

---

rollout/	
ep_len_mean	471
ep_rew_mean	-0.0891
time/	
episodes	16
fps	43
time_elapsed	174
total_timesteps	7536
train/	
actor_loss	-41.8
critic_loss	0.147
ent_coef	0.108
ent_coef_loss	-22
learning_rate	0.0003
n_updates	7435

---

rollout/	
ep_len_mean	471
ep_rew_mean	-0.0965
time/	
episodes	20
fps	43
time_elapsed	218
total_timesteps	9420
train/	
actor_loss	-38.7
critic_loss	1.01
ent_coef	0.0614
ent_coef_loss	-27.6

learning_rate	0.0003
n_updates	9319
<hr/>	
rollout/	
ep_len_mean	471
ep_rew_mean	-0.101
time/	
episodes	24
fps	43
time_elapsed	260
total_timesteps	11304
train/	
actor_loss	-36.3
critic_loss	0.975
ent_coef	0.035
ent_coef_loss	-32.7
learning_rate	0.0003
n_updates	11203
<hr/>	
rollout/	
ep_len_mean	471
ep_rew_mean	-0.102
time/	
episodes	28
fps	43
time_elapsed	304
total_timesteps	13188
train/	
actor_loss	-33.5
critic_loss	0.00989
ent_coef	0.0199
ent_coef_loss	-38.6
learning_rate	0.0003
n_updates	13087
<hr/>	
rollout/	
ep_len_mean	471
ep_rew_mean	-0.0938

time/		
episodes	32	
fps	43	
time_elapsed	348	
total_timesteps	15072	
train/		
actor_loss	-30.3	
critic_loss	0.162	
ent_coef	0.0113	
ent_coef_loss	-43.4	
learning_rate	0.0003	
n_updates	14971	

---

rollout/		
ep_len_mean	471	
ep_rew_mean	-0.0814	
time/		
episodes	36	
fps	43	
time_elapsed	392	
total_timesteps	16956	
train/		
actor_loss	-27.6	
critic_loss	0.00887	
ent_coef	0.00648	
ent_coef_loss	-47.5	
learning_rate	0.0003	
n_updates	16855	

---

rollout/		
ep_len_mean	471	
ep_rew_mean	-0.0624	
time/		
episodes	40	
fps	43	
time_elapsed	436	
total_timesteps	18840	
train/		
actor_loss	-24.3	

```
| critic_loss | 0.00392 |
| ent_coef     | 0.00371 |
| ent_coef_loss | -51.2 |
| learning_rate | 0.0003 |
| n_updates     | 18739 |

-----
Saved model to ./sb3_logs_wf\sac_window_0.zip
Saved VecNormalize to ./sb3_logs_wf\sac_vecnorm_window_0.pkl
Window 0 测试结束:
    终值资产: 12143.97 (初始 9971.90)
    总增长倍数: 1.2178
    年化收益: 70.57%
    年化波动: 11.21%
    Sharpe: 6.293
```

```
=====
Window 1
Train: 2022-07-03 ~ 2024-07-02
Test : 2024-07-03 ~ 2025-01-02
Using cuda device
```

```
C:\Users\HP\AppData\Local\Temp\ipykernel_9704\1757598842.py:64: FutureWarning: DataFrame.fillna with 'method' is deprecated and
will raise in a future version. Use obj.ffill() or obj.bfill() instead.
    feature_data.fillna(method="ffill", inplace=True)
C:\Users\HP\AppData\Local\Temp\ipykernel_9704\1757598842.py:65: FutureWarning: DataFrame.fillna with 'method' is deprecated and
will raise in a future version. Use obj.ffill() or obj.bfill() instead.
    feature_data.fillna(method="bfill", inplace=True)
C:\Users\HP\AppData\Local\Temp\ipykernel_9704\1757598842.py:64: FutureWarning: DataFrame.fillna with 'method' is deprecated and
will raise in a future version. Use obj.ffill() or obj.bfill() instead.
    feature_data.fillna(method="ffill", inplace=True)
C:\Users\HP\AppData\Local\Temp\ipykernel_9704\1757598842.py:65: FutureWarning: DataFrame.fillna with 'method' is deprecated and
will raise in a future version. Use obj.ffill() or obj.bfill() instead.
    feature_data.fillna(method="bfill", inplace=True)
```

rollout/		
ep_len_mean	471	
ep_rew_mean	0.0683	
time/		
episodes	4	
fps	46	
time_elapsed	40	
total_timesteps	1884	
train/		
actor_loss	-26.1	
critic_loss	1.37	
ent_coef	0.586	
ent_coef_loss	-5.38	
learning_rate	0.0003	
n_updates	1783	

rollout/		
ep_len_mean	471	
ep_rew_mean	0.0846	
time/		
episodes	8	
fps	45	
time_elapsed	83	
total_timesteps	3768	
train/		
actor_loss	-37.7	
critic_loss	0.35	
ent_coef	0.333	
ent_coef_loss	-11	
learning_rate	0.0003	
n_updates	3667	

rollout/		
ep_len_mean	471	
ep_rew_mean	0.0698	
time/		
episodes	12	
fps	44	

time_elapsed	126
total_timesteps	5652
train/	
actor_loss	-41.4
critic_loss	0.058
ent_coef	0.19
ent_coef_loss	-16.6
learning_rate	0.0003
n_updates	5551

---

rollout/	
ep_len_mean	471
ep_rew_mean	0.0768
time/	
episodes	16
fps	44
time_elapsed	168
total_timesteps	7536
train/	
actor_loss	-41.4
critic_loss	0.172
ent_coef	0.108
ent_coef_loss	-22
learning_rate	0.0003
n_updates	7435

---

rollout/	
ep_len_mean	471
ep_rew_mean	0.066
time/	
episodes	20
fps	44
time_elapsed	211
total_timesteps	9420
train/	
actor_loss	-38.5
critic_loss	1.08
ent_coef	0.0614
ent_coef_loss	-27.5

learning_rate	0.0003
n_updates	9319
<hr/>	
rollout/	
ep_len_mean	471
ep_rew_mean	0.0553
time/	
episodes	24
fps	44
time_elapsed	253
total_timesteps	11304
train/	
actor_loss	-36
critic_loss	1.05
ent_coef	0.035
ent_coef_loss	-32.8
learning_rate	0.0003
n_updates	11203
<hr/>	
rollout/	
ep_len_mean	471
ep_rew_mean	0.0531
time/	
episodes	28
fps	44
time_elapsed	298
total_timesteps	13188
train/	
actor_loss	-33.4
critic_loss	0.0109
ent_coef	0.0199
ent_coef_loss	-38.8
learning_rate	0.0003
n_updates	13087
<hr/>	
rollout/	
ep_len_mean	471
ep_rew_mean	0.0569

time/		
episodes	32	
fps	43	
time_elapsed	343	
total_timesteps	15072	
train/		
actor_loss	-30.1	
critic_loss	0.188	
ent_coef	0.0113	
ent_coef_loss	-43.5	
learning_rate	0.0003	
n_updates	14971	

---

rollout/		
ep_len_mean	471	
ep_rew_mean	0.0667	
time/		
episodes	36	
fps	43	
time_elapsed	385	
total_timesteps	16956	
train/		
actor_loss	-27.5	
critic_loss	0.0116	
ent_coef	0.00647	
ent_coef_loss	-48.2	
learning_rate	0.0003	
n_updates	16855	

---

rollout/		
ep_len_mean	471	
ep_rew_mean	0.0788	
time/		
episodes	40	
fps	43	
time_elapsed	429	
total_timesteps	18840	
train/		
actor_loss	-24.3	

critic_loss	0.00214
ent_coef	0.00368
ent_coef_loss	-52.8
learning_rate	0.0003
n_updates	18739

```
-----  
Saved model to ./sb3_logs_wf\sac_window_1.zip  
Saved VecNormalize to ./sb3_logs_wf\sac_vecnorm_window_1.pkl
```

Window 1 测试结束:

```
    终值资产: 11412.39 (初始 10124.11)  
    总增长倍数: 1.1272  
    年化收益: 37.40%  
    年化波动: 12.15%  
    Sharpe: 3.077
```

```
=====
```

Window 2

Train: 2023-01-03 ~ 2025-01-02

Test : 2025-01-03 ~ 2025-07-02

Using cuda device

```
C:\Users\HP\AppData\Local\Temp\ipykernel_9704\1757598842.py:64: FutureWarning: DataFrame.fillna with 'method' is deprecated and  
will raise in a future version. Use obj.ffill() or obj.bfill() instead.
```

```
    feature_data.fillna(method="ffill", inplace=True)
```

```
C:\Users\HP\AppData\Local\Temp\ipykernel_9704\1757598842.py:65: FutureWarning: DataFrame.fillna with 'method' is deprecated and  
will raise in a future version. Use obj.ffill() or obj.bfill() instead.
```

```
    feature_data.fillna(method="bfill", inplace=True)
```

```
C:\Users\HP\AppData\Local\Temp\ipykernel_9704\1757598842.py:64: FutureWarning: DataFrame.fillna with 'method' is deprecated and  
will raise in a future version. Use obj.ffill() or obj.bfill() instead.
```

```
    feature_data.fillna(method="ffill", inplace=True)
```

```
C:\Users\HP\AppData\Local\Temp\ipykernel_9704\1757598842.py:65: FutureWarning: DataFrame.fillna with 'method' is deprecated and  
will raise in a future version. Use obj.ffill() or obj.bfill() instead.
```

```
    feature_data.fillna(method="bfill", inplace=True)
```

rollout/		
ep_len_mean	472	
ep_rew_mean	0.12	
time/		
episodes	4	
fps	45	
time_elapsed	41	
total_timesteps	1888	
train/		
actor_loss	-26.7	
critic_loss	0.199	
ent_coef	0.585	
ent_coef_loss	-5.42	
learning_rate	0.0003	
n_updates	1787	

rollout/		
ep_len_mean	472	
ep_rew_mean	0.139	
time/		
episodes	8	
fps	44	
time_elapsed	84	
total_timesteps	3776	
train/		
actor_loss	-37.6	
critic_loss	0.166	
ent_coef	0.332	
ent_coef_loss	-11.1	
learning_rate	0.0003	
n_updates	3675	

rollout/		
ep_len_mean	472	
ep_rew_mean	0.13	
time/		
episodes	12	
fps	44	

time_elapsed	127
total_timesteps	5664
train/	
actor_loss	-41.2
critic_loss	5.02
ent_coef	0.189
ent_coef_loss	-16.5
learning_rate	0.0003
n_updates	5563

---

rollout/	
ep_len_mean	472
ep_rew_mean	0.133
time/	
episodes	16
fps	44
time_elapsed	170
total_timesteps	7552
train/	
actor_loss	-42.1
critic_loss	0.0771
ent_coef	0.107
ent_coef_loss	-22.3
learning_rate	0.0003
n_updates	7451

---

rollout/	
ep_len_mean	472
ep_rew_mean	0.129
time/	
episodes	20
fps	44
time_elapsed	213
total_timesteps	9440
train/	
actor_loss	-39
critic_loss	0.0419
ent_coef	0.061
ent_coef_loss	-27.9

learning_rate	0.0003
n_updates	9339

---

rollout/	
ep_len_mean	472
ep_rew_mean	0.135
time/	
episodes	24
fps	44
time_elapsed	257
total_timesteps	11328
train/	
actor_loss	-36.2
critic_loss	0.68
ent_coef	0.0347
ent_coef_loss	-32.6
learning_rate	0.0003
n_updates	11227

---

rollout/	
ep_len_mean	472
ep_rew_mean	0.13
time/	
episodes	28
fps	44
time_elapsed	300
total_timesteps	13216
train/	
actor_loss	-33.3
critic_loss	0.0109
ent_coef	0.0197
ent_coef_loss	-38.6
learning_rate	0.0003
n_updates	13115

---

rollout/	
ep_len_mean	472
ep_rew_mean	0.128

time/		
episodes	32	
fps	43	
time_elapsed	343	
total_timesteps	15104	
train/		
actor_loss	-30.4	
critic_loss	0.158	
ent_coef	0.0112	
ent_coef_loss	-44	
learning_rate	0.0003	
n_updates	15003	

---

rollout/		
ep_len_mean	472	
ep_rew_mean	0.128	
time/		
episodes	36	
fps	43	
time_elapsed	387	
total_timesteps	16992	
train/		
actor_loss	-28	
critic_loss	0.0797	
ent_coef	0.000636	
ent_coef_loss	-50.1	
learning_rate	0.0003	
n_updates	16891	

---

rollout/		
ep_len_mean	472	
ep_rew_mean	0.136	
time/		
episodes	40	
fps	43	
time_elapsed	433	
total_timesteps	18880	
train/		
actor_loss	-24.7	

critic_loss	0.00266
ent_coef	0.00362
ent_coef_loss	-54.9
learning_rate	0.0003
n_updates	18779

-----

Saved model to ./sb3\_logs\_wf\sac\_window\_2.zip  
Saved VecNormalize to ./sb3\_logs\_wf\sac\_vecnorm\_window\_2.pkl  
Window 2 测试结束:  
 终值资产: 10603.74 (初始 10025.39)  
 总增长倍数: 1.0577  
 年化收益: 16.80%  
 年化波动: 24.45%  
 Sharpe: 0.687

```
C:\Users\HP\AppData\Local\Temp\ipykernel_9704\1757598842.py:64: FutureWarning: DataFrame.fillna with 'method' is deprecated and
will raise in a future version. Use obj.ffill() or obj.bfill() instead.
    feature_data.fillna(method="ffill", inplace=True)
C:\Users\HP\AppData\Local\Temp\ipykernel_9704\1757598842.py:65: FutureWarning: DataFrame.fillna with 'method' is deprecated and
will raise in a future version. Use obj.ffill() or obj.bfill() instead.
    feature_data.fillna(method="bfill", inplace=True)
```

Out[31]:

rw_id	train_start	train_end	test_start	test_end	test_steps	start_wealth	end_wealth	total_growth	ann_return	ann_vol	sha
0	2022-01-03	2024-01-02	2024-01-03	2024-07-02	93	9971.897342	12143.969822	1.217819	0.705697	0.112133	6.293
1	2022-07-03	2024-07-02	2024-07-03	2025-01-02	95	10124.113547	11412.386682	1.127248	0.374006	0.121533	3.077
2	2023-01-03	2025-01-02	2025-01-03	2025-07-02	91	10025.393147	10603.735845	1.057688	0.168023	0.244493	0.687

## 修改部分

添加了科技股票、国债、标普500和黄金等数据作为环境输入特征

收益严重偏大，进行修改：

```
==== Portfolio Performance (deterministic rollout) ====
Periods per year (inferred): 252.00 Start wealth: 10297.47 End wealth: 23731074.08
Total growth: 2304.553868 CAGR: 804.485651% Ann. Return (geom from equity): 804.485651% Ann. Vol: 22.710210%
Sharpe (rf=0): 9.8502 Max Drawdown: -5.59% Calmar: 143.9561
```

## 1、对vec\_env进行了归一化处理

```
==== Portfolio Performance (deterministic rollout) ====
Periods per year (inferred): 252.00 Start wealth: 10180.25 End wealth: 1013360.39
Total growth: 99.541782 CAGR: 270.067928% Ann. Return (geom from equity): 270.067928% Ann. Vol: 22.074986%
Sharpe (rf=0): 6.0514 Max Drawdown: -8.50% Calmar: 31.7818
```

## 2、调整奖励函数

用协方差矩阵估计组合波动 添加每次交易成本千一

```
==== Portfolio Performance (deterministic rollout) ====
Periods per year (inferred): 252.00 Start wealth: 10055.45 End wealth: 54941.09
Total growth: 5.463811 CAGR: 62.092061% Ann. Return (geom from equity): 62.092061% Ann. Vol: 17.712724% Sharpe (rf=0): 2.8174 Max Drawdown: -11.18% Calmar: 5.5563
```

## 3、更换为sac算法

```
==== Portfolio Performance (deterministic rollout) ====
Periods per year (inferred): 252.00 Start wealth: 10145.52 End wealth: 22875.43
Total growth: 2.254731 CAGR: 26.016893% Ann. Return (geom from equity): 26.016893% Ann. Vol: 14.523819% Sharpe (rf=0): 1.6653 Max Drawdown: -12.48% Calmar: 2.0840
```

删除了step函数中对动作的softmax处理

```
==== Portfolio Performance (deterministic rollout) ====
Periods per year (inferred): 252.00 Start wealth: 10096.53 End wealth: 36203.07
Total growth: 3.585696 CAGR: 43.791829% Ann. Return (geom from equity): 43.791829% Ann. Vol: 12.005115% Sharpe (rf=0): 3.0874 Max Drawdown: -8.52% Calmar: 5.1418
```

# 优化奖励函数权重

benchmark:

```
lambda_risk = 1 lambda_turnover = 0.001 reward = base_reward - lambda_risk * risk_penalty - lambda_turnover * cost
```

```
==== Portfolio Performance (deterministic rollout) ====
Periods per year (inferred): 252.00
Start wealth: 10164.10
End wealth: 39643.65
Total growth: 3.900359
CAGR: 47.273482%
Ann. Return (geom from equity): 47.273482%
Ann. Vol: 17.140595%
Sharpe (rf=0): 2.3455
Max Drawdown: -13.67%
Calmar: 3.4585
```